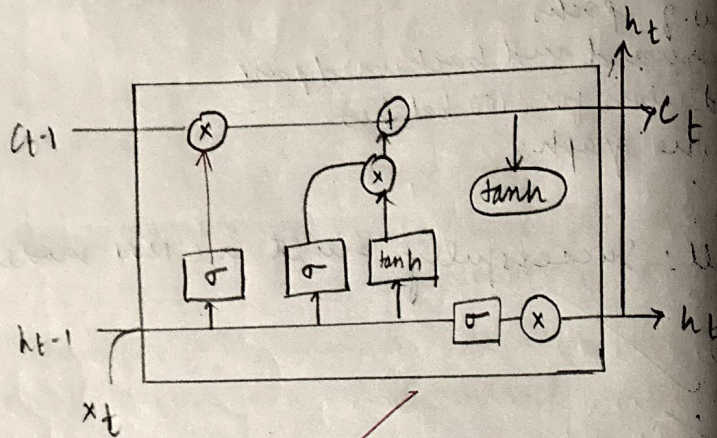


LSTM Architecture



Experiment 8

23/9/25

Build AN LSTM Model

Aim: To build and train a Long Short Term Memory (LSTM) model using PyTorch for time series forecasting

Objectives:

- To generate or load a numerical time series dataset.
- To process data and create sequential input samples for the LSTM
- To evaluate model performance on unseen test data

Pseudocode:

Start

Import necessary libraries

Generate or LOAD the time series data

Preprocess the data

- Define sequence-length
- For each i in range($\text{len}(\text{data}) - \text{sequence-length}$)
 $x[i] = \text{data}[i:i + \text{sequence-length}]$
 $y[i] = \text{data}[i + \text{sequence-length}]$
- split data into train and test set

Define the LSTM model

- Input size = 1
- Hidden size = 64
- Number of layers = 2
- Output layer = 1 neuron

Forward pass:

$\text{out}_t = \text{LSTM}(x)$

$\text{output} = \text{FullyConnected}(\text{out}[:, -1, :])$


→ Train (Predict, output, compute loss, update weights)

- Forecast future ~~new~~ values after testing the model
- Plot results.
- End.

Result: Successfully build LSTM model

~~1/9/2020~~

NOISY SINE WAVE TIME SERIES



epoch [1/25]: loss : 0.244074

epoch [2/25]: loss : 0.018688


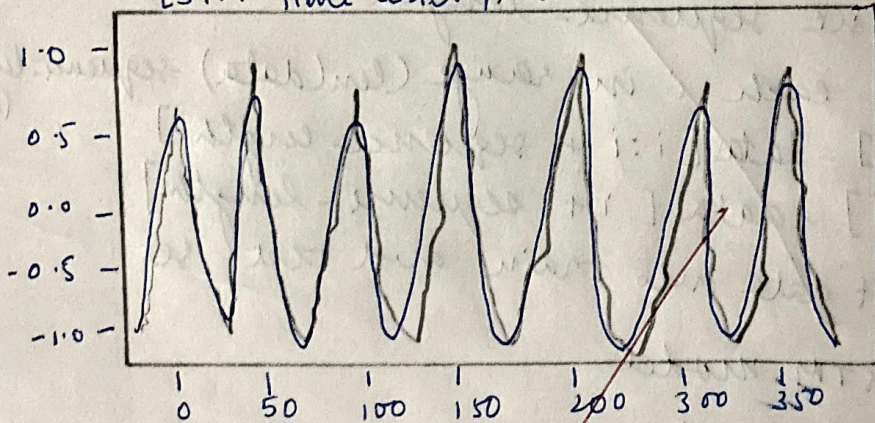
epoch [3/25]: loss : 0.014108

epoch [4/25]: loss : 0.03592

⋮

epoch [24/25]: loss 0.011999

LSTM Time series Prediction




```
In [ ]: import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
import matplotlib.pyplot as plt
from torch.utils.data import DataLoader, TensorDataset

np.random.seed(42)
time = np.arange(0, 200, 0.1)
data = np.sin(time) + 0.1 * np.random.randn(len(time))

plt.figure(figsize=(8,3))
plt.plot(time, data)
plt.title("Noisy Sine Wave Time Series")
plt.show()

def create_sequences(data, seq_length):
    xs, ys = [], []
    for i in range(len(data) - seq_length):
        x = data[i:i+seq_length]
        y = data[i+seq_length]
        xs.append(x)
        ys.append(y)
    return np.array(xs), np.array(ys)

seq_length = 50
X, y = create_sequences(data, seq_length)

X = torch.tensor(X, dtype=torch.float32).unsqueeze(-1) # (samples, seq_len, 1)
y = torch.tensor(y, dtype=torch.float32).unsqueeze(-1) # (samples, 1)
```

10:10 PM ✓

```
train_size = int(0.8 * len(X))
X_train, X_test = X[:train_size], X[train_size:]
y_train, y_test = y[:train_size], y[train_size:]

train_loader = DataLoader(TensorDataset(X_train, y_train), batch_size=32, shuffle=True)
```

```
class LSTMRegressor(nn.Module):
    def __init__(self, input_size=1, hidden_size=64, num_layers=2):
        super(LSTMRegressor, self).__init__()
        self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first=True)
        self.fc = nn.Linear(hidden_size, 1)

    def forward(self, x):
        out, _ = self.lstm(x)
        out = out[:, -1, :] # use last time step
        out = self.fc(out)
        return out
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = LSTMRegressor().to(device)
criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
num_epochs = 25
for epoch in range(num_epochs):
    model.train()
    total_loss = 0
    for batch_X, batch_y in train_loader:
        batch_X, batch_y = batch_X.to(device), batch_y.to(device)
        optimizer.zero_grad()
        output = model(batch_X)
        loss = criterion(output, batch_y)
        loss.backward()
```

10:10 PM ✓

```

total_loss += loss.item()
print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {total_loss/len(train_loader):.6f}")

model.eval()
with torch.no_grad():
    preds = model(X_test.to(device)).cpu().numpy()
    actual = y_test.cpu().numpy()

plt.figure(figsize=(8,4))
plt.plot(range(len(actual)), actual, label='Actual')
plt.plot(range(len(preds)), preds, label='Predicted')
plt.legend()
plt.title("LSTM Time Series Prediction")
plt.show()


with torch.no_grad():
    seq = X_test[-1].unsqueeze(0).to(device)
    future_preds = []
    for _ in range(50):
        pred = model(seq)
        future_preds.append(pred.item())
        seq = torch.cat([seq[:, 1:, :], pred.unsqueeze(1)], dim=1)

plt.figure(figsize=(8,3))
plt.plot(range(len(data)), data, label='Original Data')
plt.plot(range(len(data), len(data)+50), future_preds, label='Future Prediction')
plt.legend()
plt.title("Future Forecasting with LSTM")
plt.show()

```

10:11 PM ✓

Noisy Sine Wave Time Series




```

Epoch [1/25], Loss: 0.244074
Epoch [2/25], Loss: 0.018688
Epoch [3/25], Loss: 0.014108
Epoch [4/25], Loss: 0.013900
Epoch [5/25], Loss: 0.013592
Epoch [6/25], Loss: 0.012818
Epoch [7/25], Loss: 0.012583
Epoch [8/25], Loss: 0.012952
Epoch [9/25], Loss: 0.013004
Epoch [10/25], Loss: 0.013327
Epoch [11/25], Loss: 0.013180
Epoch [12/25], Loss: 0.012638
Epoch [13/25], Loss: 0.012094
Epoch [14/25], Loss: 0.011980
Epoch [15/25], Loss: 0.012476
Epoch [16/25], Loss: 0.012440
Epoch [17/25], Loss: 0.011978
Epoch [18/25], Loss: 0.011622


```

10:11 PM ✓

LSTM Time Series Prediction



Future Forecasting with LSTM



Aim: To develop and train a Recurrent Neural Network

Objective:

Generate sequential data

Build an RNN model

Train the model

Visualize predicted vs actual values

Pseudocode:

Start

Import required libraries

Set hyperparameters

$TIMESTEPS = 10$

$RNN_UNITS = 32$

$EPOCHS = 100$

$BATCH_SIZE = 16$

Generate synthetic sequential data

Create dataset for RNN:

Reshape input data

Split data into training & testing data

Build RNN Model:

model = Sequential()

Add SimpleRNN layer

Add Dense layer

Train the model


history = model.fit(x_train, y_train,
epochs=EPOCHS, validation_split=0.1)

Evaluate model on test data

loss, mae = model.evaluate(x_test, y_test)

rmse = sqrt(loss)

RNN structure



predict using trained model
visualize predictions:
end.

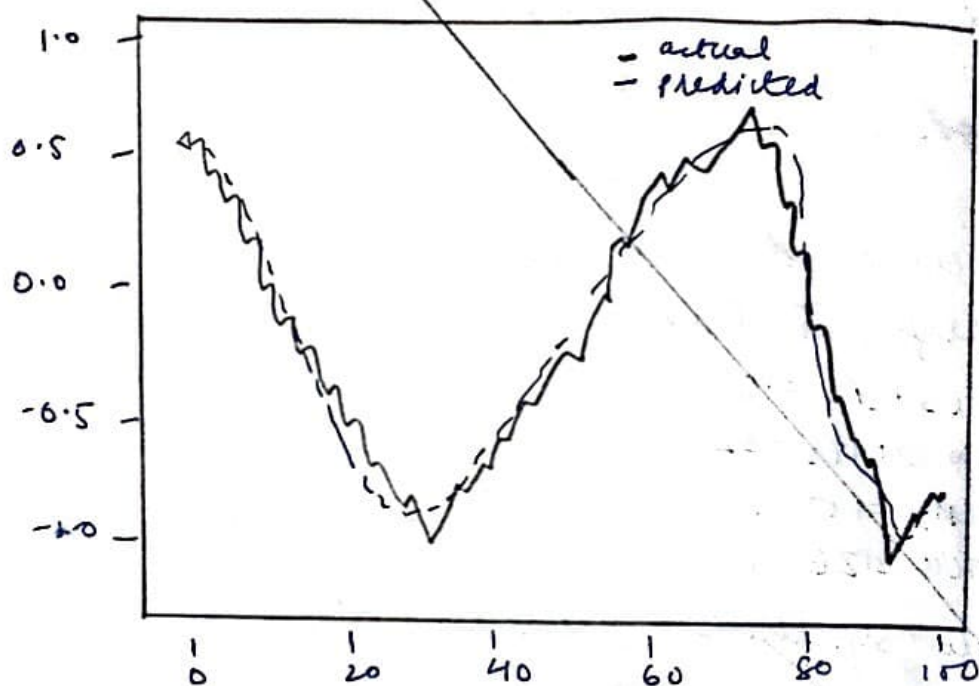
result: Successfully built RNN model.

~~11/19/2018~~


Test Loss (MSE) = 0.0169

Test MAE : 0.1070

Test RMSE : 0.1299



AIRLINE PASSENGERS FORECASTING



```

import torch
import torch.nn as nn
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from torch.utils.data import Dataset, DataLoader
import seaborn as sns

```

```

data = sns.load_dataset("flights")
all_data = data['passengers'].values.astype(float)

```

```

test_data_size = 12
train_data = all_data[:-test_data_size]
test_data = all_data[-test_data_size:]

```

```

scaler = MinMaxScaler(feature_range=(-1, 1))
train_data_normalized = scaler.fit_transform(train_data.reshape(-1, 1))
train_data_normalized = torch.FloatTensor(train_data_normalized).view(-1)

```

```

def create_inout_sequences(input_data, tw):
    inout_seq = []
    L = len(input_data)
    for i in range(L - tw):
        train_seq = input_data[i:i + tw]
        train_label = input_data[i + tw:i + tw + 1]
        inout_seq.append((train_seq, train_label))
    return inout_seq

```

```

train_window = 12
train_inout_seq = create_inout_sequences(train_data_normalized, train_window)

```

```

class TimeSeriesDataset(Dataset):
    def __init__(self, sequences):

```

10:14 PM ✓

```

class TimeSeriesDataset(Dataset):
    def __init__(self, sequences):
        super().__init__()
        self.sequences = sequences
    def __len__(self):
        return len(self.sequences)
    def __getitem__(self, idx):
        return self.sequences[idx][0], self.sequences[idx][1]

```

```

train_dataset = TimeSeriesDataset(train_inout_seq)
batch_size = 10
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)

```

```

class SimpleRNN(nn.Module):
    def __init__(self, input_size=1, hidden_layer_size=100, output_size=1):
        super().__init__()
        self.hidden_layer_size = hidden_layer_size

        self.rnn = nn.RNN(input_size, hidden_layer_size, batch_first=True)
        self.linear = nn.Linear(hidden_layer_size, output_size)

    def forward(self, input_seq):
        rnn_input = input_seq.unsqueeze(-1)
        h0 = torch.zeros(1, rnn_input.size(0), self.hidden_layer_size).to(input_seq.device)
        rnn_out, h_n = self.rnn(rnn_input, h0)
        predictions = self.linear(rnn_out[:, -1, :])
        return predictions

```

```
input_dim = 1
```

10:14 PM ✓


```

optimizer.zero_grad()

y_pred = model(seq)

single_loss = loss_function(y_pred, labels)

single_loss.backward()
optimizer.step()

if epoch % 20 == 0:
    print(f'Epoch {epoch:3} Loss: {single_loss.item():10.8f}')

print(f'Final Loss: {single_loss.item():10.8f}')
print("Training complete!")

model.eval()

test_input = scaler.transform(test_data.reshape(-1, 1))
test_input = torch.FloatTensor(test_input).view(-1)

fut_pred = 12
test_inputs = train_data_normalized[-train_window:].tolist()

for i in range(fut_pred):
    seq = torch.FloatTensor(test_inputs[-train_window:])

    with torch.no_grad():

        y_pred = model(seq.unsqueeze(0)).squeeze()

        test_inputs.append(y_pred.item())

actual_predictions = test_inputs[-fut_pred:]

```

10:15 PM ✓


Starting training SimpleRNN on 120 sequences...

```

Epoch 0 Loss: 0.11418664
Epoch 20 Loss: 0.00629161
Epoch 40 Loss: 0.00531148
Epoch 60 Loss: 0.01036970
Epoch 80 Loss: 0.00373696
Final Loss: 0.00587350
Training complete!

```

Airline Passengers Forecasting with Simple PyTorch RNN



10:15 PM ✓