

Experiment 4 BUILD A SIMPLE FEED FORWARD
14.8.2025 NEURAL NETWORK TO RECOGNIZE HANDWRITTEN
CHARACTER

Aim: To design and implement a simple feedforward neural network that can classify handwritten characters

Objective: To design and implement a FNN including forward propagation, activation functions and backpropagation for training and understand it

Pseudocode:

- Load the Dataset
 - Normalize pixel values
- Initialize the Network
 - Define input layer size
 - Define hidden layer size
 - Define output layer size
 - Initialize weights and biases randomly.
- Define activation functions
- Do forward and backward propagation
- Training loop:
 - For number of epochs:
 - Forward propagate all training samples
 - Backpropagate errors
 - update weights and biases
- Evaluate Model:
- Output.

Notes:

- Device configuration is done so that code can automatically decide whether to run on CPU or GPU.
- MNIST images have pixel values from 0 to 255 so normalization rescales pixels to a smaller range
- MNIST contains 60,000 training data and 10,000 testing images
- Input layer (784 neurons = 28×28 pixels) $h_1 = W_1 u + b_1$
Hidden layer (128 neurons + ReLU activation) $\Rightarrow h_2 = W_2 a_1 + b_2$
Output layer (10 neurons with softmax) $\Rightarrow z = W_3 a_2 + b_3$
(logit)
- logit refers to the raw, unnormalized outputs from the final layer of a model before any activation function is applied.
- Input image \rightarrow Flatten - fc1
Apply ReLU \rightarrow non-linear activation
fc2 \rightarrow Apply ReLU
fc3 \rightarrow get logits
Pass logits to cross entropy \rightarrow compare with the true label to compute loss

Observation

Epoch [1/5]: Loss: 0.4134

Epoch [2/5]: Loss: 0.1975

Epoch [3/5], Loss: 0.1445

Epoch [4/5], Loss: 0.1153

Epoch [5/5], Loss: 0.0977

Test accuracy: 97.06%


```

import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms

# 1. Device configuration
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# 2. Transform: normalize pixel values
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

# 3. Load MNIST dataset
train_dataset = torchvision.datasets.MNIST(root="./data", train=True, download=True, transform=transform)
test_dataset = torchvision.datasets.MNIST(root="./data", train=False, download=True, transform=transform)

train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=64, shuffle=False)

# 4. Build a simple Feed Forward Neural Network
class FFN(nn.Module):
    def __init__(self):
        super(FFN, self).__init__()
        self.fc1 = nn.Linear(28*28, 128) # input -> hidden
        self.fc2 = nn.Linear(128, 64)    # hidden -> hidden
        self.fc3 = nn.Linear(64, 10)     # hidden -> output

    def forward(self, x):
        x = x.view(-1, 28*28) # flatten image
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))

```

```
model = FFN().to(device)
```

```
# 5. Loss and optimizer
```

```
criterion = nn.CrossEntropyLoss()
```

```
c
```

```
# 6. Training Loop
```

```
epochs = 5
```

```
for epoch in range(epochs):
```

```
    running_loss = 0.0
```

```
    for images, labels in train_loader:
```

```
        images, labels = images.to(device), labels.to(device)
```

```
        optimizer.zero_grad()
```

```
        outputs = model(images)
```

```
        loss = criterion(outputs, labels)
```

```
        loss.backward()
```

```
        optimizer.step()
```

```
        running_loss += loss.item()
```

```
    print(f"Epoch [{epoch+1}/{epochs}], Loss: {running_loss/len(train_loader):.4f}")
```

```
# 7. Testing Loop
```

```
correct = 0
```

```
total = 0
```

```
with torch.no_grad():
```

```
    for images, labels in test_loader:
```

```
        images, labels = images.to(device), labels.to(device)
```

```
        outputs = model(images)
```

```
        _, predicted = torch.max(outputs, 1)
```

```
        total += labels.size(0)
```

```
        correct += (predicted == labels).sum().item()
```



```
Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82
Attempting uninstall: nvidia-cuda-nvrtc-cu12
Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82
Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:
  Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82
Attempting uninstall: nvidia-cuda-cupti-cu12
Found existing installation: nvidia-cuda-cupti-cu12 12.5.82
Uninstalling nvidia-cuda-cupti-cu12-12.5.82:
  Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82
Attempting uninstall: nvidia-cublas-cu12
Found existing installation: nvidia-cublas-cu12 12.5.3.2
Uninstalling nvidia-cublas-cu12-12.5.3.2:
  Successfully uninstalled nvidia-cublas-cu12-12.5.3.2
Attempting uninstall: nvidia-cusparse-cu12
Found existing installation: nvidia-cusparse-cu12 12.5.1.3
Uninstalling nvidia-cusparse-cu12-12.5.1.3:
  Successfully uninstalled nvidia-cusparse-cu12-12.5.1.3
Attempting uninstall: nvidia-cudnn-cu12
Found existing installation: nvidia-cudnn-cu12 9.3.0.75
Uninstalling nvidia-cudnn-cu12-9.3.0.75:
  Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75
Attempting uninstall: nvidia-cusolver-cu12
Found existing installation: nvidia-cusolver-cu12 11.6.3.83
Uninstalling nvidia-cusolver-cu12-11.6.3.83:
  Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Successfully installed nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-cu12-12.4.127 nvidia-cudnn-cu12-9.1.0.70 nvidia-cufft-cu12-11.2.1.3 nvidia-curand-cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-cusparse-cu12-12.3.1.170 nvidia-nccl-cu12-2.21.5 nvidia-nvjitlink-cu12-12.4.127
Epoch [1/5], Loss: 0.4134
Epoch [2/5], Loss: 0.1975
Epoch [3/5], Loss: 0.1445
Epoch [4/5], Loss: 0.1153
Epoch [5/5], Loss: 0.0977
Test Accuracy: 97.06%
```

```
import matplotlib.pyplot as plt
import torchvision

# Get one batch of images
examples = iter(train_loader)
images, labels = next(examples)

# Plot first 6 images
plt.figure(figsize=(8, 4))
for i in range(6):
    plt.subplot(2, 3, i+1)
    plt.imshow(images[i][0], cmap='gray') # images[i][0] because MNIST is grayscale
    plt.title(f"Label: {labels[i].item()}")
    plt.axis('off')
plt.show()
```

Label: 2



Label: 0



Label: 5



Label: 0



Label: 3



Label: 7



```
plt.plot(epochs, loss_values, marker='o')  
plt.xlabel("Epoch")  
plt.ylabel("Loss")  
plt.title("Training Loss Curve")  
plt.grid(True)  
plt.show()
```

