

# Experiments 1/9/2025 Activation Functions & Their roles

Aim: To learn about different activation functions and their roles.

Objective:

- To evaluate and compare the performance of different activation functions
- By training each model for a fixed number of epochs and measuring its test accuracy on unseen MNIST test data.

Pseudocode:

- Load MNIST dataset
  - Define neural network architecture
  - Define list of activation function:  
Sigmoid, Tanh, ReLU, LeakyReLU, SiLU
  - Train model for 5 epochs
    - Forward pass
    - Compute loss
    - Backward pass
    - Update weights
  - Compute accuracy
  - Display result.
- END

Observation

- Overall high accuracy as MNIST digit recognition task is relatively simple and all activation functions handle it well.



→ Among the activations, ReLU achieved the best accuracy because it avoids vanishing gradients and trains faster.

Sigmoid - 97.74%.

Tanh - 97.24%.

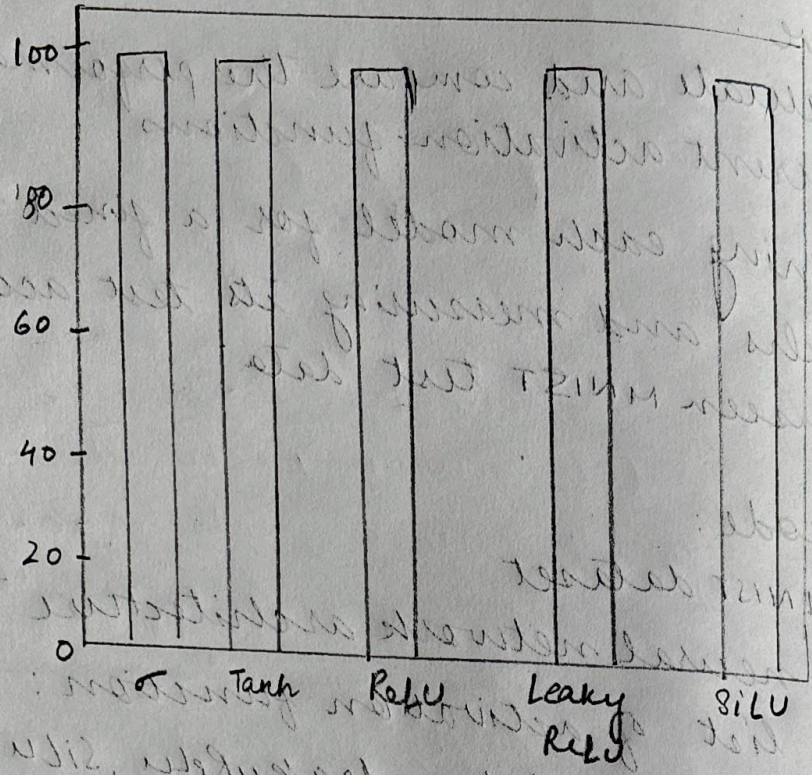
ReLU - 97.85%.

Leaky ReLU : 97.84%.

SiLU : 97.62%.

Result: Trained the model and learned about different activation functions successfully.





- Forward pass
- Compute loss
- Backward pass
- Update weights

will high accuracy on MNIST digit recognition task is relatively simple and all activation functions handle it well.

```
import torch
import torch.nn as nn
import torch.nn.functional as F

x = torch.tensor([-2.0, -1.0, 0.0, 1.0, 2.0])

# Using module
sigmoid = nn.Sigmoid()
print(sigmoid(x))

# Using functional API
print(F.sigmoid(x))  # Deprecated, better to use torch.sigmoid
print(torch.sigmoid(x))

tensor([0.1192, 0.2689, 0.5000, 0.7311, 0.8808])
tensor([0.1192, 0.2689, 0.5000, 0.7311, 0.8808])
tensor([0.1192, 0.2689, 0.5000, 0.7311, 0.8808])

tanh = nn.Tanh()
print(tanh(x))
print(torch.tanh(x))  # Functional

tensor([-0.9640, -0.7616,  0.0000,  0.7616,  0.9640])
tensor([-0.9640, -0.7616,  0.0000,  0.7616,  0.9640])

relu = nn.ReLU()
print(relu(x))
print(F.relu(x))
```



```
In [ ]: leaky_relu = nn.LeakyReLU(negative_slope=0.01)
        print(leaky_relu(x))
        print(F.leaky_relu(x, negative_slope=0.01))

        tensor([-0.0200, -0.0100,  0.0000,  1.0000,  2.0000])
        tensor([-0.0200, -0.0100,  0.0000,  1.0000,  2.0000])
```

```
In [ ]: elu = nn.ELU()
        print(elu(x))
        print(F.elu(x))

        tensor([-0.8647, -0.6321,  0.0000,  1.0000,  2.0000])
        tensor([-0.8647, -0.6321,  0.0000,  1.0000,  2.0000])
```

```
In [ ]: y = torch.tensor([2.0, 1.0, 0.1])

        softmax = nn.Softmax(dim=0)
        print(softmax(y))
        print(F.softmax(y, dim=0))

        tensor([0.6590, 0.2424, 0.0986])
        tensor([0.6590, 0.2424, 0.0986])
```

```
In [ ]: import torch
        import torch.nn.functional as F
        import matplotlib.pyplot as plt

        # Input values
        x = torch.linspace(-5, 5, 200)
```

```
# Activation functions
```

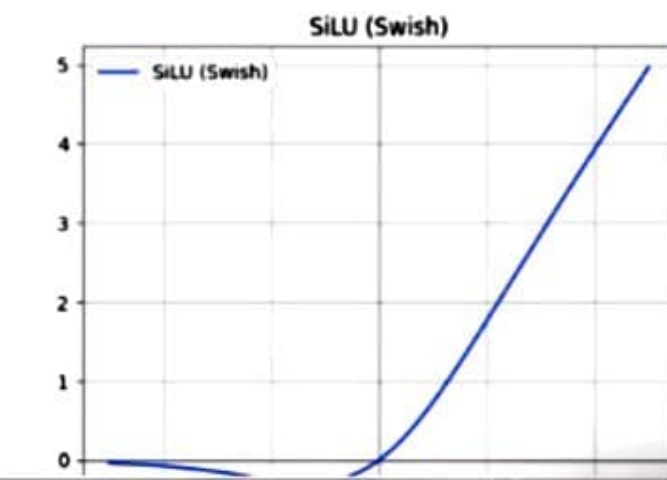
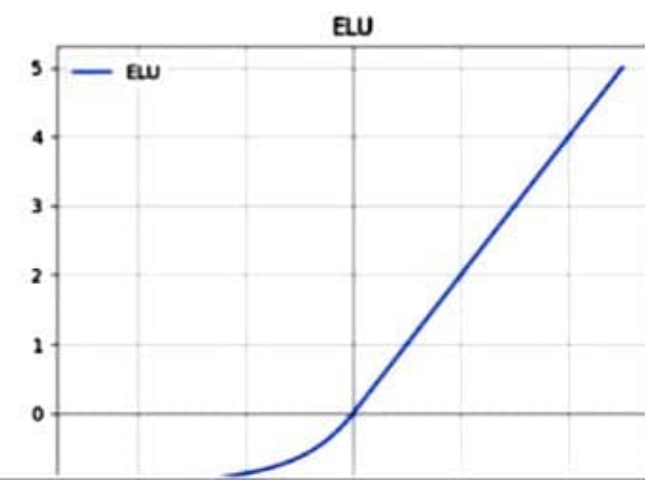
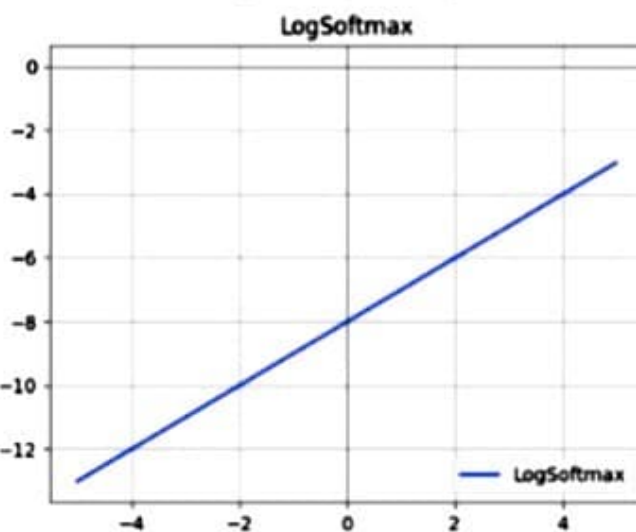
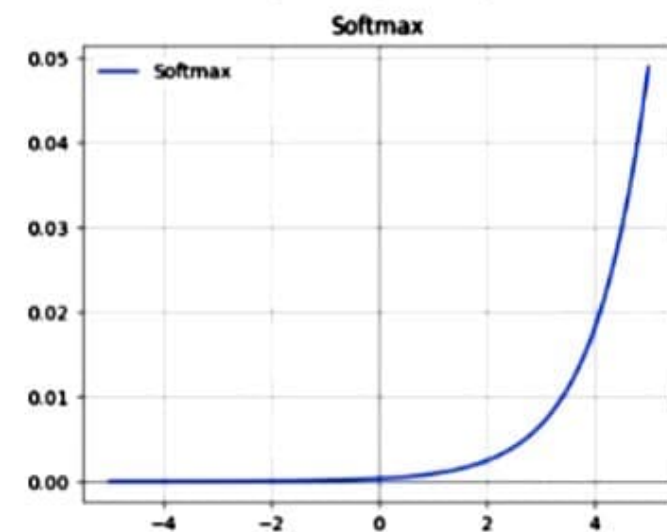
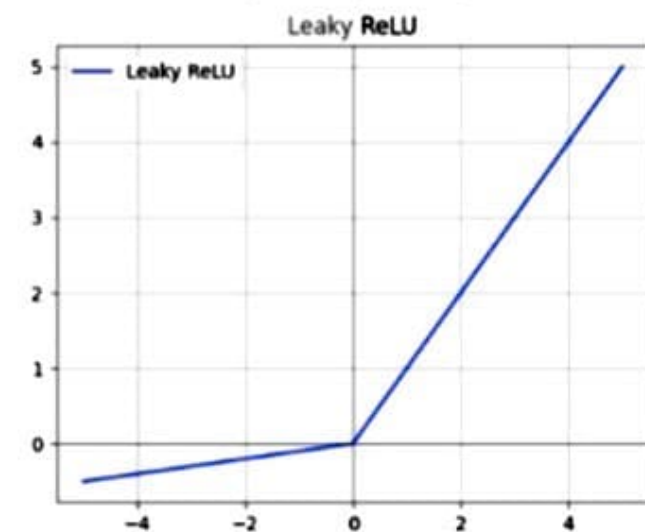
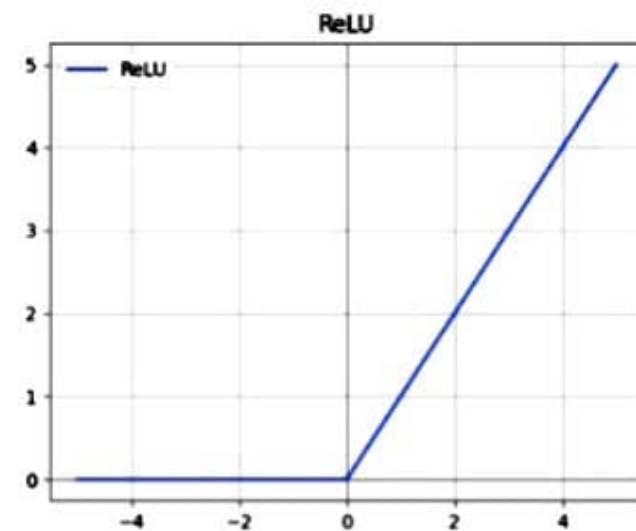

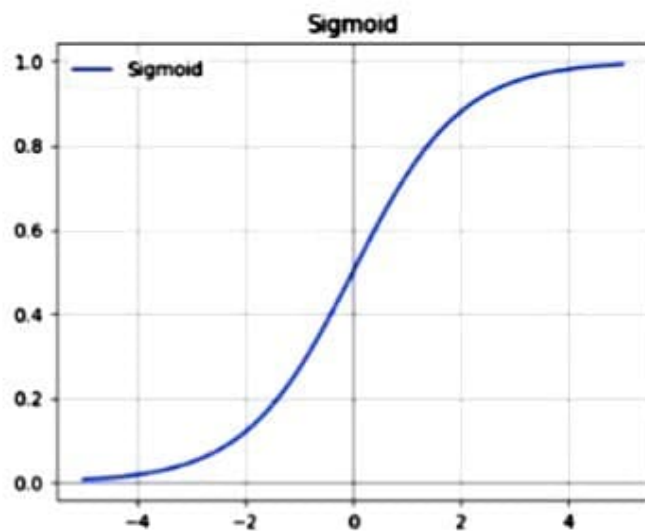
```
activations = {  
    "Sigmoid": torch.sigmoid(x),  
    "Tanh": torch.tanh(x),  
    "ReLU": F.relu(x),  
    "Leaky ReLU": F.leaky_relu(x, negative_slope=0.1),  
    "Softmax": F.softmax(x, dim=0),  
    "LogSoftmax": F.log_softmax(x, dim=0),  
    "ELU": F.elu(x, alpha=1.0),  
    "SiLU (Swish)": F.silu(x)  
}
```

```
# Plotting
```

```
plt.figure(figsize=(15, 12))
```

```
for i, (name, y) in enumerate(activations.items(), 1):  
    plt.subplot(3, 3, i)  
    plt.plot(x.numpy(), y.detach().numpy(), label=name, color="blue")  
    plt.title(name)  
    plt.axhline(0, color="black", linewidth=0.5)  
    plt.axvline(0, color="black", linewidth=0.5)  
    plt.grid(True)  
    plt.legend()
```

```
plt.tight_layout()  
plt.show()
```



```

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt

# Device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Data
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

train_dataset = datasets.MNIST(root="./data", train=True, transform=transform, download=True)
test_dataset = datasets.MNIST(root="./data", train=False, transform=transform, download=True)

train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=1000, shuffle=False)

# Model with configurable activation
class SimpleNN(nn.Module):
    def __init__(self, activation_fn):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(28*28, 256)
        self.fc2 = nn.Linear(256, 128)
        self.fc3 = nn.Linear(128, 10)
        self.activation = activation_fn

    def forward(self, x):

```



100%		9.91M/9.91M	[00:00<00:00, 33.4MB/s]
100%		28.9k/28.9k	[00:00<00:00, 1.06MB/s]
100%		1.65M/1.65M	[00:00<00:00, 9.38MB/s]
100%		4.54k/4.54k	[00:00<00:00, 7.18MB/s]

Training with Sigmoid activation...

Test Accuracy (Sigmoid): 97.74%

Training with Tanh activation...

Test Accuracy (Tanh): 97.24%

Training with ReLU activation...

Test Accuracy (ReLU): 97.85%

Training with LeakyReLU activation...

Test Accuracy (LeakyReLU): 97.45%

Training with SiLU activation...

Test Accuracy (SiLU): 97.62%

=== Final Comparison ===

Sigmoid: 97.74%

Tanh: 97.24%


ReLU: 97.85%

LeakyReLU: 97.45%

SiLU: 97.62%

4:12 PM ✓

#### Activation Function Comparison on MNIST



4:12 PM ✓