

Experiment 6

Implement Gradient Descent

9/9/2025

& Backpropagation

Aim: To implement a deep neural network demonstrating the use of forward Propagation, backpropagation and gradient descent.

Objectives:

- Understand the XOR problem as a non linearly separable dataset.
- Perform backpropagation to compute gradients.
- Visualize the training loss and decision boundary of the network.

Pseudocode:

- Import PyTorch
- Define XOR dataset
- Initialize hidden layer
- Initialize output layer
- Use sigmoid
- Define the forward function
- Initialize Model:

model = XORNet()

loss_function = MSE Loss()

- Set number of epochs
 - Train the model
 - End.
- For each epoch
- Forward pass
 - compute loss
 - Backpropagation: loss.backward
 - update weights
 - Reset gradients

Pseudocode: Plotting Decision Boundary

- Import libraries
- create a grid of points covering input space
- Plot decision boundary using contours
- blue = 0, red = 1
- end.

Observation

epoch 0, Loss: 0.2595

epoch 1000, Loss: 0.2497

:

epoch 9000, Loss: 0.0194

:

epoch 19000 / Loss: 0.0024

tensor ([[0.0494],
[0.9517],
[0.9560],
[0.455]])

loss is computed using $L = \frac{1}{N} \sum_{i=1}^N (\overset{\text{Predicted}}{\hat{y}_i} - \overset{\text{true}}{y_i})^2$

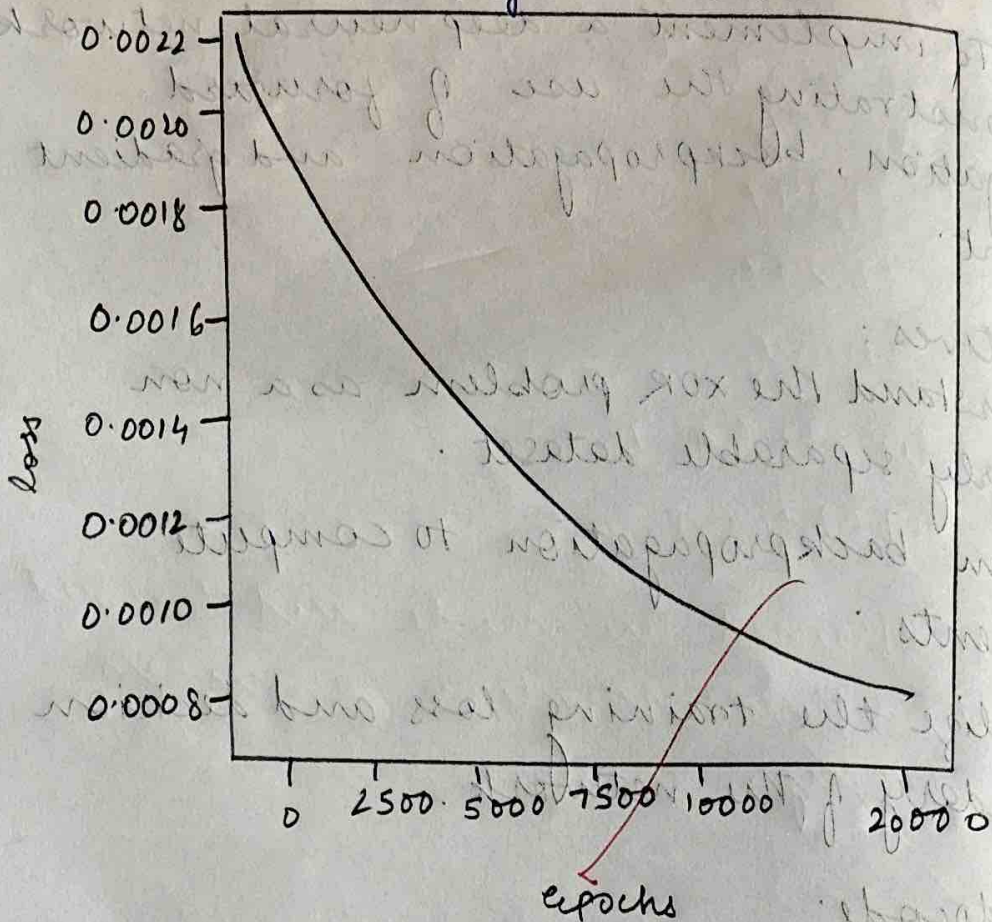
gradients using

$$\frac{\partial L}{\partial w} = (\hat{y} - y) \cdot \hat{y} \cdot (1 - \hat{y}) \cdot a_{\text{prev}}$$

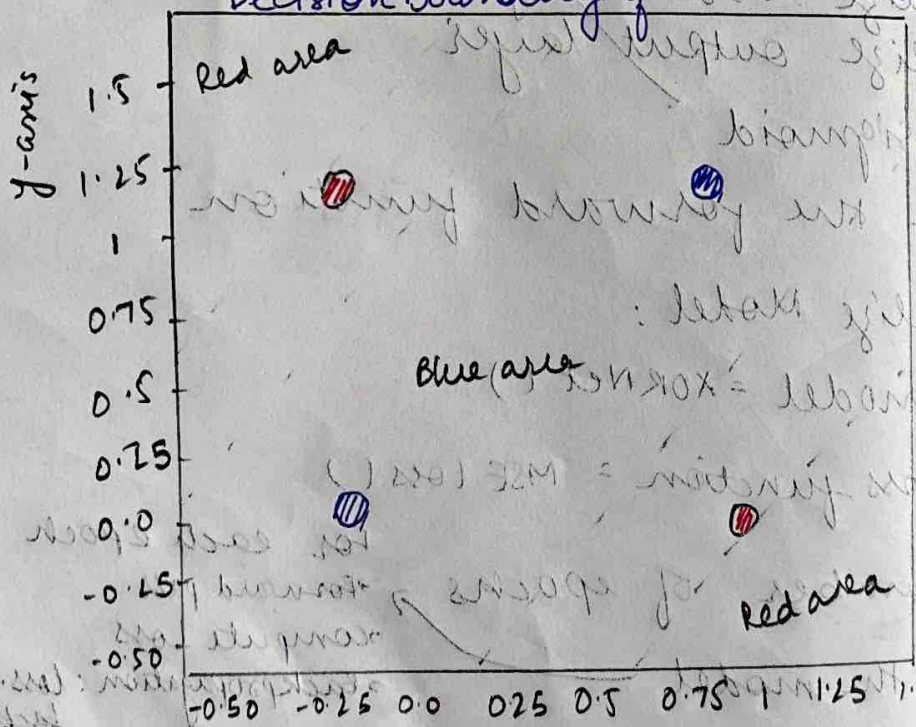
↓
input to
this
weight.

$$W = W - \eta \frac{\partial L}{\partial W}$$

Training Loss Curve



Decision Boundary of Neural Network



Result: Successfully implemented backpropagation to train a deep neural network.

```

import torch
import torch.nn as nn
import torch.optim as optim

# Dataset (XOR problem)
X = torch.tensor([[0,0],[0,1],[1,0],[1,1]], dtype=torch.float32)
y = torch.tensor([[0],[1],[1],[0]], dtype=torch.float32)

# Define a simple neural network
class XORNet(nn.Module):
    def __init__(self):
        super(XORNet, self).__init__()
        self.hidden = nn.Linear(2, 4) # input layer -> hidden
        self.output = nn.Linear(4, 1) # hidden -> output
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        x = self.sigmoid(self.hidden(x))
        x = self.sigmoid(self.output(x))
        return x

# Initialize model, loss, and optimizer
model = XORNet()
criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.1)

# Training Loop
epochs = 20000
for epoch in range(epochs):
    # Forward pass
    outputs = model(X)
    loss = criterion(outputs, y)

```

5:09 PM ✓

```

    # Backward pass + update
    optimizer.zero_grad() # reset gradients
    loss.backward()       # compute gradients
    optimizer.step()       # update weights

    if epoch % 1000 == 0:
        print(f"Epoch {epoch}, Loss: {loss.item():.4f}")

# Final predictions
print("Final predictions:")
print(model(X).detach())

```

```

Epoch 0, Loss: 0.2595
Epoch 1000, Loss: 0.2497
Epoch 2000, Loss: 0.2492
Epoch 3000, Loss: 0.2483
Epoch 4000, Loss: 0.2458
Epoch 5000, Loss: 0.2361
Epoch 6000, Loss: 0.1899
Epoch 7000, Loss: 0.0877
Epoch 8000, Loss: 0.0361
Epoch 9000, Loss: 0.0194
Epoch 10000, Loss: 0.0125
Epoch 11000, Loss: 0.0089
Epoch 12000, Loss: 0.0069
Epoch 13000, Loss: 0.0055
Epoch 14000, Loss: 0.0046
Epoch 15000, Loss: 0.0039
Epoch 16000, Loss: 0.0034
Epoch 17000, Loss: 0.0030
Epoch 18000, Loss: 0.0027
Epoch 19000, Loss: 0.0024
Final predictions:
tensor([[0.0494],
        [0.9517],
        [0.9569],
        [0.0431]])


```

5:10 PM ✓

```
Epoch 19000, Loss: 0.0024  
Final predictions:  
tensor([[0.0494],  
        [0.9517],  
        [0.9560],  
        [0.0455]])
```

```
In [6]: import matplotlib.pyplot as plt  
  
losses = []  
  
for epoch in range(epochs):  
    outputs = model(X)  
    loss = criterion(outputs, y)  
  
    optimizer.zero_grad()  
    loss.backward()  
    optimizer.step()  
  
    losses.append(loss.item())  
  
# Plot Loss curve  
plt.plot(losses)  
plt.xlabel("Epochs")  
plt.ylabel("Loss")  
plt.title("Training Loss Curve")  
plt.show()
```

5:10 PM ✓



```
In [7]: import numpy as np
```

5:10 PM ✓


```
xx, yy = np.meshgrid(np.linspace(-0.5, 1.5, 200),  
                    np.linspace(-0.5, 1.5, 200))  
grid = torch.tensor(np.c_[xx.ravel(), yy.ravel()], dtype=torch.float32)  
  
# Get predictions  
with torch.no_grad():  
    Z = model(grid).numpy()  
Z = Z.reshape(xx.shape)  
  
# Plot decision boundary  
plt.contourf(xx, yy, Z, levels=50, cmap="coolwarm", alpha=0.7)  
plt.scatter(X[:,0], X[:,1], c=y[:,0], edgecolor='k', s=100, cmap="coolwarm")  
plt.title("Decision Boundary of Neural Network")  
plt.show()
```

Decision Boundary of Neural Network

