

# Experiment 14 Implement a Pre Trained CNN Model as a feature extractor using transfer learning

Aim: To perform image classification on CIFAR-10 dataset using Transfer learning with a pre-trained ResNet 18 model.

## Objectives

→ To understand and implement transfer learning using a pre-trained ResNet 18 model.

→ To train the modified model and visualize the loss curve

## Pseudo code.

→ Begin

→ Import torch, torchvision, torch.nn

→ set device

→ load pre-trained ResNet 18 model:

→ Freeze all layers:

→ for each parameter in model  
    set requires\_grad = False

→ modify the final layer for CIFAR-10:

trainset ← CIFAR10(train=True)

→ Define transformation:

transform ← Resize(224 × 224) + ToTensor()



## Observation

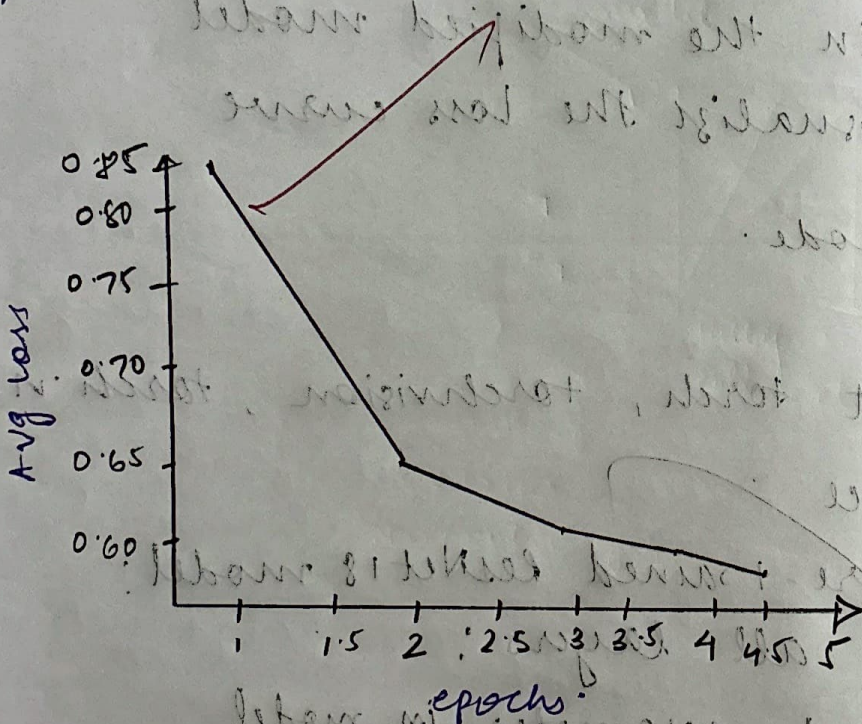
epoch 1: loss : 0.8348

epoch 2: loss : 0.6289

epoch 3: loss : 0.5937

epoch 4: loss : 0.5997

epoch 5: loss : 0.5699



loss curve



define loss function and optimizer:

criterion  $\leftarrow$  crossentropyloss()

for each epoch in range(1 to 10):

avg loss  $\leftarrow$  running loss / total-batches

Append avg-loss to loss-values

Print epoch number and avg loss

print "Training Complete"

plots loss curve:

x-axis  $\leftarrow$  epochs

y-axis  $\leftarrow$  average loss

end.

Result: Successfully implemented a pretrained CNN model.

Q7

```
import matplotlib.pyplot as plt
```

```
# Device setup
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
print("Using device:", device)
```

```
# Load pre-trained ResNet18
```

```
model = models.resnet18(pretrained=True)
```

```
# Freeze all convolutional layers to use as feature extractor
```

```
for param in model.parameters():
```

```
    param.requires_grad = False
```

```
# Replace the final layer to match CIFAR-10 (10 classes)
```

```
model.fc = nn.Linear(model.fc.in_features, 10)
```

```
model = model.to(device)
```

```
# Data transformations
```

```
transform = transforms.Compose([
```

```
    transforms.Resize((224, 224)),
```

```
    transforms.ToTensor()
```

```
])
```

```
# Load dataset (CIFAR-10)
```

```
trainset = CIFAR10(root='./data', train=True, download=True, transform=transform)
```

```
trainloader = DataLoader(trainset, batch_size=64, shuffle=True)
```

```
# Loss and optimizer
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = optim.Adam(model.fc.parameters(), lr=0.001)
```

```
# Lists to store loss values for graph
```

```
loss_values = []
```

```
# Training for a few epochs
```



```

# Lists to store loss values for graph
loss_values = []

# Training for a few epochs
epochs = 5
for epoch in range(epochs):
    running_loss = 0.0
    for images, labels in trainloader:
        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    avg_loss = running_loss / len(trainloader)
    loss_values.append(avg_loss)
    print(f"Epoch [{epoch+1}/{epochs}], Loss: {avg_loss:.4f}")

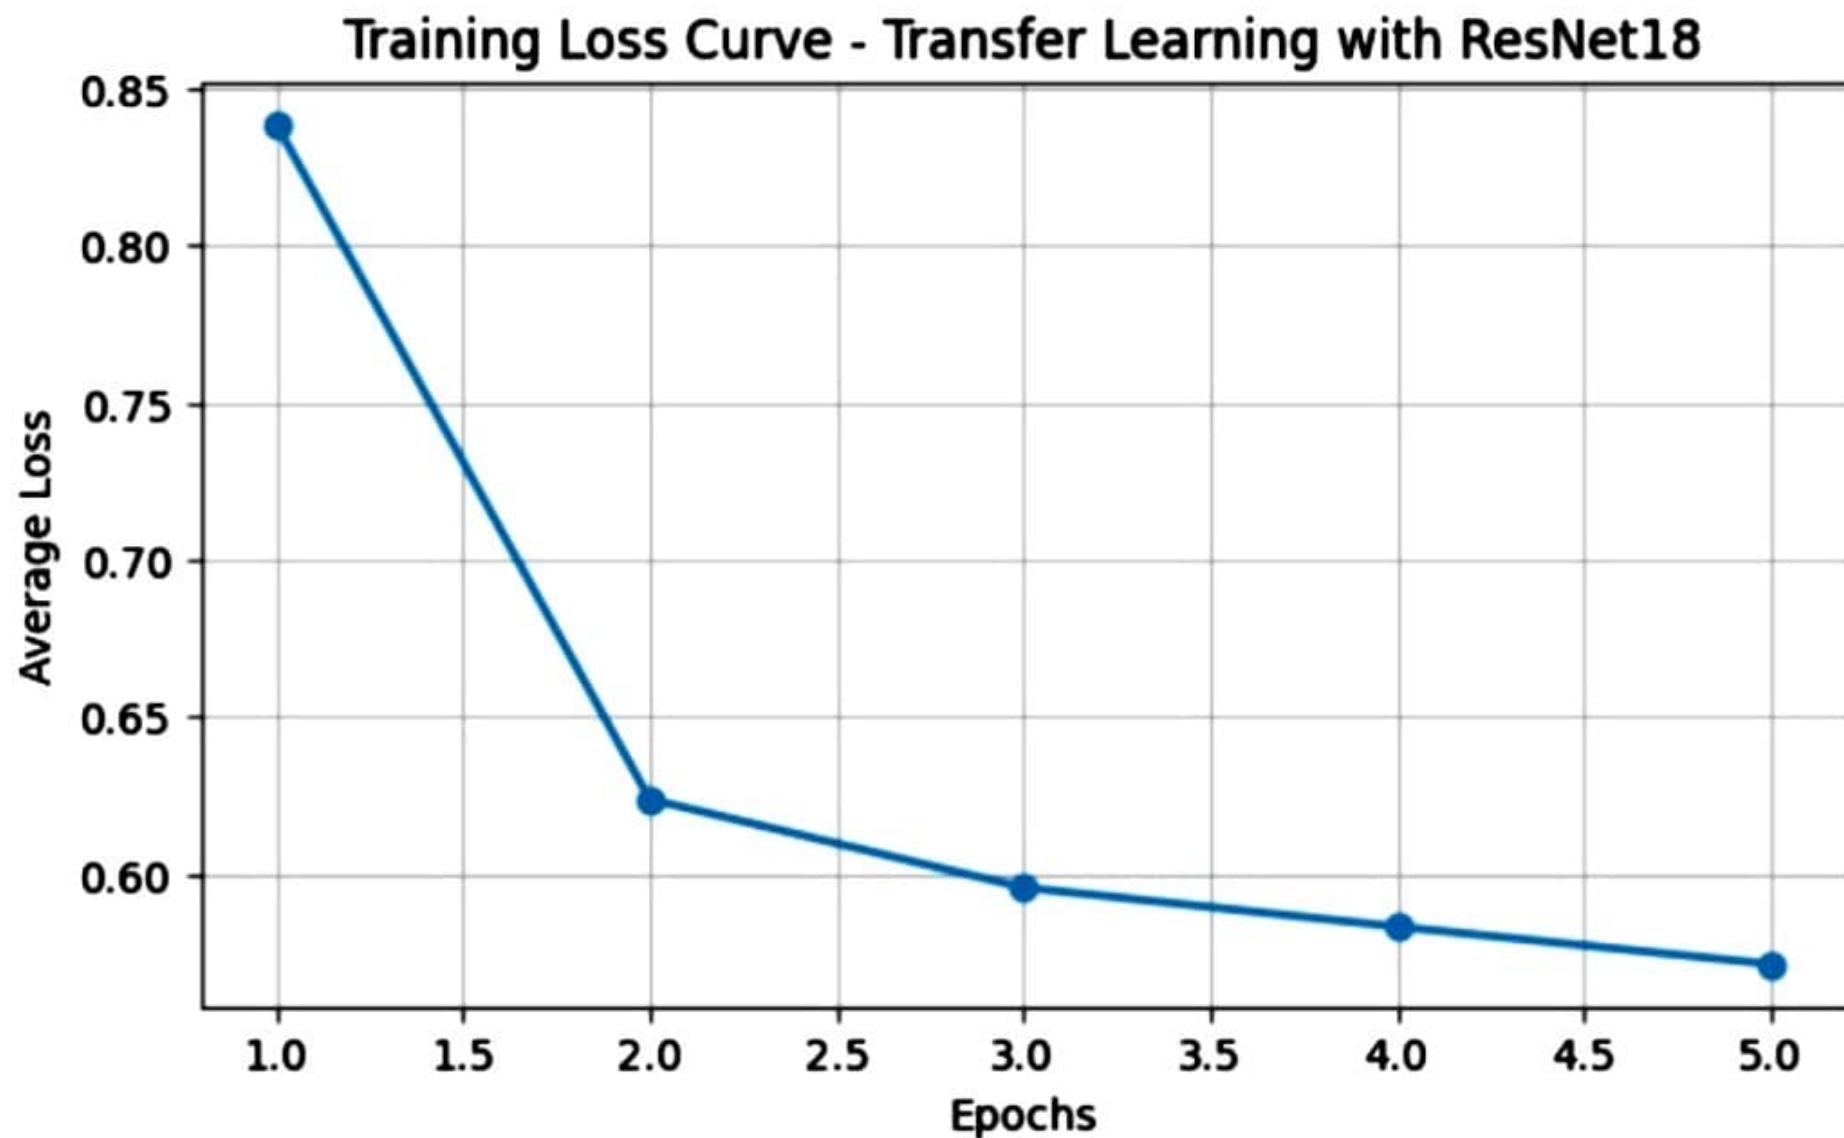
print("\n✅ Training complete using Pre-trained ResNet18 as Feature Extractor.\n")

# Plot the training Loss graph
plt.figure(figsize=(7, 4))
plt.plot(range(1, epochs + 1), loss_values, marker='o', linestyle='-', linewidth=2)
plt.title("Training Loss Curve - Transfer Learning with ResNet18")
plt.xlabel("Epochs")
plt.ylabel("Average Loss")
plt.grid(True)
plt.show()

```

Epoch [1/5], Loss: 0.8388  
Epoch [2/5], Loss: 0.6238  
Epoch [3/5], Loss: 0.5957  
Epoch [4/5], Loss: 0.5832  
Epoch [5/5], Loss: 0.5713

✅ Training complete using Pre-trained ResNet18 as Feature Extractor.





Experiment 15  
27/10/25

Implement a YOLO  
model to detect objects

Aim: Demonstrate how to perform  
object detection on sample using a  
pre-trained CNN model

Objective

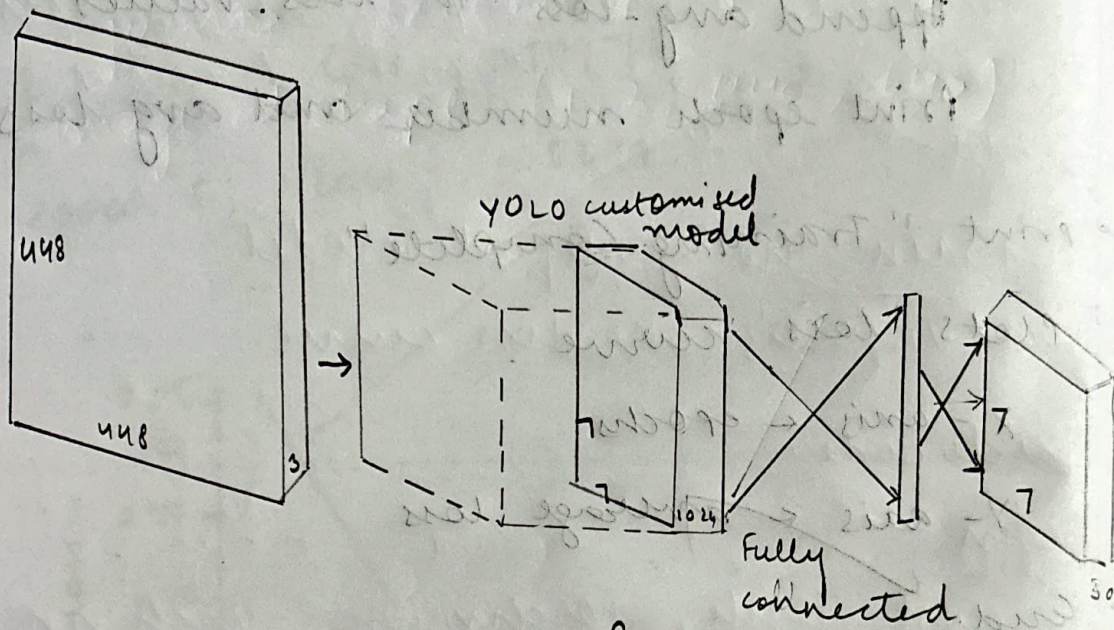
→ loading and understanding the  
YOLO model that was trained on  
COCO dataset

Pseudocode

- Import necessary libraries YOLO  
from ultralytics
- `SELMODEL = load_yolo_model("yolov8n.pt")`
- load the image as a URL form
- Run prediction
  - SET confidence - threshold = 0.7
  - SET results - list = `model(image-pil,`  
`conf = confidence - threshold)`
  - SET result = `results - list[0]`
- Draw detection
  - || Use built-in `plot()` method to  
draw boxes and labels
  - || return a numpy array in BGR format



# YOLO Model Architecture



consists of

Conv layer

~~Maxpool layer~~



output

448 x 640 1car, 322.2ms

Speed  $\approx$  34.0 ms preprocess, 322.2ms

inference, 32.6 ms postprocess per

image at shape (1, 3, 448, 640)

Load the image on a new form  
model = load\_model('model.h5')

prediction  
T = threshold = 0.7  
conf = model.predict(img)  
conf = conf - threshold  
[0] 123 - results = results

Draw detection  
It will print the plot of the image  
with boxes and labels  
It returns a numpy array in the format

convert from BGR to RGB from  
matplotlib

Initialize plot and display plot.

end.

Result: Successfully detected images using  
a YOLO model

Eg. 1



```
# 1. Install dependencies
# We only need 'ultralytics' (which includes torch) and 'matplotlib'
!pip install ultralytics matplotlib --quiet

# 2. Import libraries
from ultralytics import YOLO
from PIL import Image
import matplotlib.pyplot as plt
import requests
import io # Needed for handling image from requests

# 3. Load pretrained YOLO model
# We'll use YOLOv8n (nano), which is fast and COCO-pretrained.
model = YOLO('yolov8n.pt')

# 4. Load a sample image from URL
url = "https://images.unsplash.com/photo-1503376780353-7e6692767b70" # has cars
img = None
try:
    # Get the image data
    response = requests.get(url)
    response.raise_for_status() # Raise an error for bad responses

    # Open the image from the in-memory bytes
    img = Image.open(io.BytesIO(response.content)).convert("RGB")

except Exception as e:
    print(f"Could not load image from URL. Error: {e}")
    # Create a dummy image if loading fails
    img = Image.new('RGB', (800, 600), color='gray')

# 5. Run Prediction
# YOLO handles all preprocessing (transforms, ToTensor) automatically.
# We set 'conf=0.7' to match the original code's confidence threshold.
```

Creating new Ultralytics Settings v0.0.6 file 

View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.json'

Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs\_dir=path/to/dir'. For help see <https://docs.ultralytics.com/quickstart/#ultralytics-settings>.

Downloading <https://github.com/ultralytics/assets/releases/download/v8.3.0/yolov8n.pt> to 'yolov8n.pt': 100% — 6.2MB 79.5MB/s 0.1s

0: 448x640 1 car, 322.2ms

Speed: 34.0ms preprocess, 322.2ms inference, 32.6ms postprocess per image at shape (1, 3, 448, 640)

### YOLOv8 Object Detection (COCO Pretrained)

